
Database Systems 2021 - Assignment 2

By team Shalabi: Hasan Khadra (leader),
Mahmood Darwish, Mohamad Dwik,
Mohammad Shahin

Introduction

Migrating DVD rental database from PostgreSQL to DynamoDB.



DynamoDB

- Fast and flexible NoSQL database service for any scale
- Managed and developed by Amazon.
- Designing the database tables depends purely on the type of the queries.
- Pros: It is simple to set-up. Best suited for key-value type of operations. Unstructured data.
- Cons: Won't work particularly great for relational operations. Weak querying model. Harder to predict costs.



Amazon DynamoDB

Local Installation

- Create AWS account at <https://aws.amazon.com/dynamodb/>.
- Download DynamoDB .tar file specific to your area from the given website.
- After you download the archive, extract the contents, and copy the extracted directory to a location of your choice.
- To start DynamoDB on your computer, open a command prompt window, navigate to the directory where you extracted DynamoDBLocal.jar, and enter the following command:

```
java -Djava.library.path=./DynamoDBLocal_lib -jar DynamoDBLocal.jar -sharedDb
```



Amazon DynamoDB

Migration

- 1) Understand the structure of the data.
- 2) Know the purpose of the migration.
- 3) Choose the best way to migrate the data.



Amazon DynamoDB

Migration - Our way

- Python API - PGSQL (psycopg2) - DynamoDB (boto3)
- Retrieve table names, attributes names, and attribute types
- Retrieve the actual data
- Process the data (attribute types, data)
- Push the data to DynamoDB



Amazon DynamoDB

Migration - performance


- Downloading the data from PG ~ 0.2 seconds.
- Processing the data ~ 2 seconds
- Uploading the data to Dynamodb ~ 6 minutes 23 seconds.
- Total time ~ 6 minutes 25 seconds.



Amazon DynamoDB

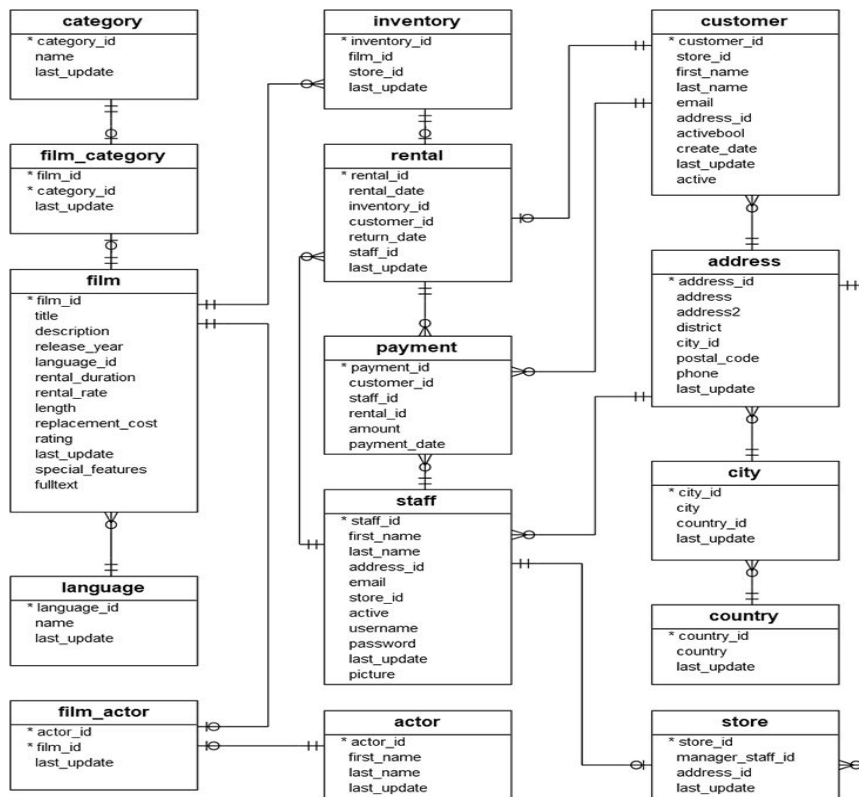
Continuous Integration

- 1) Triggered integration.
- 2) Timed integration.
- 3) Synchronize data in both DBs.



Amazon DynamoDB

Postgres UML Diagram



DynamoDB Tables Design

Design differs between RDBMS and NoSQL.

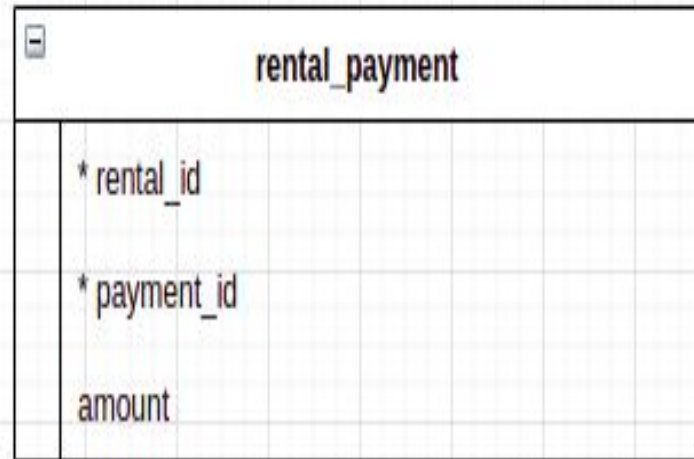








- RDBMS design
- DynamoDB design.

DynamoDB Tables Design

The tables from Postgres were moved as they are.

New table was added to handle query 1.

DynamoDB UML Diagram



Query Analysis

Query 1:

- Postgres syntax:

```
SELECT *, (SELECT COUNT(*) FROM rental r2, payment p2
WHERE r2.rental_id = p2.rental_id AND
p2.amount < p.amount) AS count_smaller_pay FROM rental
r, payment p WHERE r.rental_id = p.rental_id;
```

- Time taken to run the query: 31.84 seconds.

Query Analysis

Query 1:

- With DynamoDB, no join operation.
- We designed an additional table to make reading operation possible.



The diagram shows a table named `rental_payment` on a grid background. The table has a composite primary key with two attributes: `* rental_id` and `* payment_id`, both marked with an asterisk. A third attribute, `amount`, is listed below the key attributes.

rental_payment	
* rental_id	
* payment_id	
amount	

Query Analysis

Query 1:

- DynamoDB syntax
- Time taken to run the query 5 minutes 2 seconds.

```
amount = table.query(  
    KeyConditionExpression=Key('rental_id').eq(rental_id)  
    & Key('payment_id').eq(payment_id)  
)['Items'][0]['amount']
```

```
response = table.scan(  
    Select='COUNT',  
    FilterExpression=Key('amount').lt(amount)  
)
```

Query Analysis

Query 2:

- INSERT an actor to the actors table.
- Postgres syntax:

```
INSERT INTO actor (first_name, last_name,  
last_update) VALUES ('Mohammad', 'Dwik',  
datetime.datetime(2021, 4, 23, 0, 3, 20, 612992));
```

- Time taken to run the query ~0.0079 seconds.

Query Analysis

Query 2:

- INSERT an actor to the actors relation.
- DynamoDB syntax
- Time taken to run the query ~0.036 seconds

```
response = client.put_item(  
    TableName='actor',  
    Item={  
        'first_name': {  
            'S': 'Mahmood'  
        },  
        'last_name': {  
            'S': 'Darwish'  
        },  
        'actor_id': {  
            'N': '204'  
        }  
    },  
    ReturnConsumedCapacity='TOTAL'  
)
```

Query Analysis

Query 2:

- DELETE an actor from the actors table.
- Postgres syntax:

```
DELETE FROM actor WHERE actor_id = 201;
```

- Time taken to run the query ~0.0056 seconds.

Query Analysis

Query 2:

- DELETE an actor from the actors table.
- DynamoDB syntax
- Time taken to run the query ~0.0089 seconds.

```
response = client.delete_item(  
    TableName='actor',  
    Key={  
        'actor_id': {'N': 201}  
    }  
)
```

Query Analysis

Query 2:

- UPDATE an actor in the actors table.
- Postgres syntax:

```
UPDATE actor SET first_name = 'Hasan', last_name =  
'Khadra, last_update = datetime.datetime(2021, 4,  
23, 0, 1, 20, 976351) WHERE actor_id = 1;
```

- Time taken to run the query ~0.0099 seconds.

Query Analysis

Query 2:

- UPDATE an actor in the actors table.
- DynamoDB syntax
- Time taken to run the query ~0.0096

```
response = client.update_item(  
    ExpressionAttributeNames={  
        '#AT': 'first_name',  
        '#Y': 'last_name',  
    },  
    ExpressionAttributeValues={  
        ':t': { 'S': 'Vlad', },  
        ':y': { 'S': 'Yalalov', },  
    },  
    Key={'actor_id': { 'N': 201, }},  
    TableName='actor',  
    UpdateExpression='SET #Y = :y, #AT = :t'  
)
```

Test Functions

- Test showing that the data was successfully migrated.
- Test showing the migration process.

Thanks for listening!